

```
In [59]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [60]: df=pd.read_csv('data.csv')
df.head()
```

```
Out[60]:
```


	Bp	Sg	Al	Su	Rbc	Bu	Sc	Sod	Pot	Hemo	Wbcc	Rbcc	Htn	Class
0	80.0	1.020	1.0	0.0	1.0	36.0	1.2	137.53	4.63	15.4	7800.0	5.20	1.0	1
1	50.0	1.020	4.0	0.0	1.0	18.0	0.8	137.53	4.63	11.3	6000.0	4.71	0.0	1
2	80.0	1.010	2.0	3.0	1.0	53.0	1.8	137.53	4.63	9.6	7500.0	4.71	0.0	1
3	70.0	1.005	4.0	0.0	1.0	56.0	3.8	111.00	2.50	11.2	6700.0	3.90	1.0	1
4	80.0	1.010	2.0	0.0	1.0	26.0	1.4	137.53	4.63	11.6	7300.0	4.60	0.0	1

```
In [61]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Bp          400 non-null    float64
1    Sg          400 non-null    float64
2    Al          400 non-null    float64
3    Su          400 non-null    float64
4    Rbc         400 non-null    float64
5    Bu          400 non-null    float64
6    Sc          400 non-null    float64
7    Sod         400 non-null    float64
8    Pot         400 non-null    float64
9    Hemo        400 non-null    float64
10   Wbcc        400 non-null    float64
11   Rbcc        400 non-null    float64
12   Htn         400 non-null    float64
13   Class       400 non-null    int64
dtypes: float64(13), int64(1)
memory usage: 43.9 KB
```


```
In [62]: sns.countplot(data=df,x="Rbc",hue="Class")
```

```
Out[62]: <Axes: xlabel='Rbc', ylabel='count'>
```




```
In [63]: sns.countplot(data=df, x="Al", hue="Class")
```

```
Out[63]: <Axes: xlabel='Al', ylabel='count'>
```




```
In [64]: sns.histplot(data=df,x="Hemo",hue="Class")
```

```
Out[64]: <Axes: xlabel='Hemo', ylabel='Count'>
```



```
In [65]: sns.histplot(data=df,x="Bp",hue="Class",multiple="stack") # multiple=stack to have
```

```
Out[65]: <Axes: xlabel='Bp', ylabel='Count'>
```



Data Preprocessing

```
In [66]: # Checking for Nan Values in the dataset  
df.isnull().sum()
```

```
Out[66]: Bp      0  
Sg      0  
Al      0  
Su      0  
Rbc     0  
Bu      0  
Sc      0  
Sod     0  
Pot     0  
Hemo    0  
Wbcc    0  
Rbcc    0  
Htn     0  
Class   0  
dtype: int64
```


```
In [67]: df = df.fillna(0)  
df.isnull().sum()
```

```
Out[67]: Bp      0
         Sg      0
         Al      0
         Su      0
         Rbc     0
         Bu      0
         Sc      0
         Sod     0
         Pot     0
         Hemo    0
         Wbcc    0
         Rbcc    0
         Htn     0
         Class   0
         dtype: int64
```

Checking for balance in the class

```
In [68]: sns.countplot(data=df, x="Class", hue="Class")
         df['Class'].value_counts()
```

```
Out[68]: Class
         1    250
         0    150
         Name: count, dtype: int64
```




```
In [69]: # Using oversampling algorithm to balance the class
         from sklearn.utils import resample
         df_majority_count=df[(df['Class']==1)]
```

```
df_minority_count=df[(df['Class']==0)  
  
df_minor_upsampled=resample(df_minority_count,n_samples=250,random_state=0)  
df=pd.concat([df_minor_upsampled,df_majority_count])
```

```
In [70]: sns.countplot(data=df,x="Class",hue="Class")  
df['Class'].value_counts()
```


```
Out[70]: Class  
0      250  
1      250  
Name: count, dtype: int64
```



Boxplot - To display outliers


```
In [71]: sns.boxplot(data=df,x="Bp")
```

```
Out[71]: <Axes: xlabel='Bp'>
```




```
In [72]: sns.boxplot(data=df, x="Hemo")
```

```
Out[72]: <Axes: xlabel='Hemo'>
```




```
In [73]: sns.boxplot(data=df,x="Rbcc")
```

```
Out[73]: <Axes: xlabel='Rbcc'>
```




```
In [74]: sns.boxplot(data=df,x="Sod")
```

```
Out[74]: <Axes: xlabel='Sod'>
```

```
In [75]: sns.boxplot(data=df, x="Pot")
```

```
Out[75]: <Axes: xlabel='Pot'>
```



Dealing with outliers using Z-score method


```
In [76]: import scipy.stats as stats
```

```
z=np.abs(stats.zscore(df))
df=df[(z<3).all(axis=1)]
df.shape
```

```
Out[76]: (420, 14)
```

```
In [77]: sns.heatmap(df.corr())
```

```
Out[77]: <Axes: >
```



Dropping RBC Column since it doesn't contain any information/data


```
In [78]: df = df.drop(columns=['Rbc'])
```

```
In [79]: df.head()
```

Out[79]:

	Bp	Sg	Al	Su	Bu	Sc	Sod	Pot	Hemo	Wbcc	Rbcc	Htn	Class
297	60.0	1.025	0.0	0.0	26.0	1.0	146.0	4.9	15.8	7700.0	5.2	0.37	0
367	60.0	1.025	0.0	0.0	41.0	1.1	139.0	3.8	17.4	6700.0	6.1	0.00	0
317	70.0	1.020	0.0	0.0	48.0	1.2	139.0	4.3	15.0	8100.0	4.9	0.00	0
353	60.0	1.020	0.0	0.0	37.0	0.6	150.0	5.0	13.6	5800.0	4.5	0.00	0
259	80.0	1.020	0.0	0.0	31.0	1.2	135.0	5.0	16.1	4300.0	5.2	0.00	0

```
In [80]: corr=df.corr()["Class"]
plt.plot(corr)
plt.xticks(rotation=90)
plt.show()
```



```
In [81]: from sklearn.model_selection import train_test_split
```

```
x=df.drop('Class',axis=1)
y=df['Class']
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=0)
```

```
In [82]: from sklearn.metrics import accuracy_score,f1_score,confusion_matrix,classification
```

```
In [83]: def model_fit(model):
          model.fit(x_train,y_train)
```

```

y_pred=model.predict(x_test)
model_acc=round(accuracy_score(y_test,y_pred)*100,2)
print(f"The accuracy of the model is: {model_acc}%")
print("")
print("f1_score: ",f1_score(y_test,y_pred))
print("precision_score: ",precision_score(y_test,y_pred))
print("recall_score: ",recall_score(y_test,y_pred))
print("")
print("Confusion Matrix: ")
sns.heatmap(confusion_matrix(y_test,y_pred),annot=True,square=True,cmap='coolwa
plt.ylabel("Actual Label")
plt.xlabel("Predicted Label")
plt.title(f"confusion matrix")

print("")
print("Classification Report: ")
print(classification_report(y_test,y_pred))

```

Applying algorithms

Logistic Regression

```

In [84]: from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
model_fit(model)

```

The accuracy of the model is: 96.19%

f1_score: 0.9523809523809523
precision_score: 0.9302325581395349
recall_score: 0.975609756097561

Confusion Matrix:

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.95	0.97	64
1	0.93	0.98	0.95	41
accuracy			0.96	105
macro avg	0.96	0.96	0.96	105
weighted avg	0.96	0.96	0.96	105

```
c:\Users\karth\Documents\CSE\3rd year\6th sem\ML\lab\venv\Lib\site-packages\sklearn  
\linear_model\_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status  
=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```


Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```



KNN

```
In [85]: from sklearn.neighbors import KNeighborsClassifier  
knn = KNeighborsClassifier(n_neighbors=7)  
model_fit(knn)
```

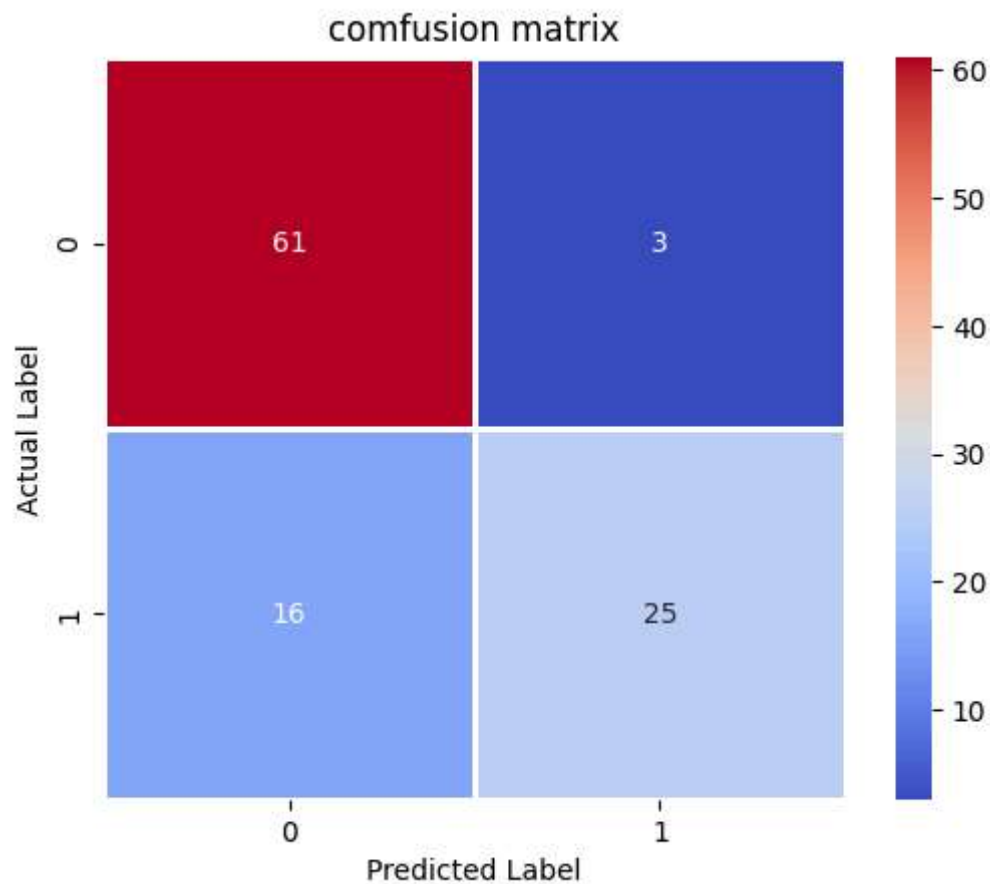
The accuracy of the model is: 81.9%

f1_score: 0.7246376811594203
 precision_score: 0.8928571428571429
 recall_score: 0.6097560975609756

Confusion Matrix:

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.95	0.87	64
1	0.89	0.61	0.72	41
accuracy			0.82	105
macro avg	0.84	0.78	0.79	105
weighted avg	0.83	0.82	0.81	105



Random Forest Algorithm

```
In [86]: from sklearn.ensemble import RandomForestClassifier
classifier = RandomForestClassifier(n_estimators=10, criterion="entropy")
model_fit(classifier)
```


The accuracy of the model is: 99.05%

f1_score: 0.9876543209876543
 precision_score: 1.0
 recall_score: 0.975609756097561

Confusion Matrix:

Classification Report:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	64
1	1.00	0.98	0.99	41
accuracy			0.99	105
macro avg	0.99	0.99	0.99	105
weighted avg	0.99	0.99	0.99	105



Ada Boost Algorithm

```
In [87]: from sklearn.ensemble import AdaBoostClassifier
abc = AdaBoostClassifier(n_estimators=50, learning_rate=1)
model_fit(abc)
```

The accuracy of the model is: 99.05%

f1_score: 0.9876543209876543

precision_score: 1.0

recall_score: 0.975609756097561

Confusion Matrix:

Classification Report:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	64
1	1.00	0.98	0.99	41
accuracy			0.99	105
macro avg	0.99	0.99	0.99	105
weighted avg	0.99	0.99	0.99	105

